



Paths in graphs

Distance in graphs

- The *distance* between two nodes is the length of the shortest path between them.
- How do we find the shortest paths from s to all other vertices?



Breadth-first search

procedure `bfs`(G, s)

Input: Graph $G = (V, E)$, directed or undirected; vertex $s \in V$

Output: For all vertices u reachable from s , `dist`(u) is set to the distance from s to u .

for all $u \in V$:

`dist`(u) = ∞

`dist`(s) = 0

$Q = [s]$ (queue containing just s)

while Q is not empty:



$u = \text{eject}(Q)$

 for all edges $(u, v) \in E$:

 if `dist`(v) = ∞ :

`inject`(Q, v)

`dist`(v) = `dist`(u) + 1



Order of visitation	Queue contents after processing node
<i>S</i>	[<i>S</i>]
<i>A</i>	[<i>A C D E</i>]
<i>C</i>	[<i>C D E B</i>]
<i>D</i>	[<i>D E B</i>]
<i>E</i>	[<i>E B</i>]
<i>B</i>	[<i>B</i>]
	[]

Correctness



- Use induction
- For each $d = 0, 1, 2, \dots$, there is a moment at which
 - (1) all nodes at distance $\leq d$ from s have their distances correctly set
 - (2) all other nodes have their distances set to ∞
 - (3) the queue contains exactly the nodes at distance d .

Analysis


- Each vertex is put on the queue exactly once $\rightarrow 2|V|$ queue operations
- **for** loop looks at each edge once (in directed graphs) or twice (in undirected graphs) $\rightarrow O(|E|)$ time
- $O(|V| + |E|)$

Weighted graphs

- Breadth-first search finds shortest paths in any graph whose edges have unit length.
- Can we adapt it to a more general graph $G = (V, E)$ whose edge lengths are *positive integers*?



- Set an alarm – estimated times of arrival
 - ex) initially, A : $T=100$, B : $T=200$
- At $T=100$, reset the alarm for B as $T=150$



Algorithm

- Set an alarm clock for node s at time 0.
- Repeat until there are no more alarms:

Say the next alarm goes off at time T , for node u . Then:

- 1) The distance from s to u is T .
- 2) For each neighbor v of u in G :
 - If there is no alarm yet for v ,
set one for time $T + l(u, v)$.
 - If v 's alarm is set for later than $T + l(u, v)$,
then reset it to this earlier time.

Priority queue

Data structure supporting the following operations :


- *Insert*. Add a new element to the set.
- *Decrease-key*. Accommodate the decrease in key value of a particular element
- *Delete-min*. Return the element with the smallest key, and remove it from the set.
- *Make-queue*. Build a priority queue out of the given elements, with the given key values. (In many implementations, this is faster than inserting the elements one by one.)

Binary heap

- Complete Binary Tree - All levels are completely filled except possibly the lowest, which is filled from the left up to a point.
- The value of each node \leq value of its children. (min-heap)





Heap data structure

Complete binary tree implemented by an array







- The root is stored in $A[1]$
- The parent of $A[i]$ is $A[\lfloor \frac{i}{2} \rfloor]$.
- The left child of $A[i]$ is $A[2 \cdot i]$.
- The right child of $A[i]$ is $A[2 \cdot i + 1]$.
- The node in the far right of the bottom level is stored in $A[n]$.
- If $2i + 1 > n$, then the node does not have a right child.

Insert



Delete-min



Dijkstra's algorithm

procedure dijkstra(G, l, s)

Input: Graph $G = (V, E)$, directed or undirected;
 positive edge lengths $\{l_e : e \in E\}$; vertex $s \in V$

Output: For all vertices u reachable from s , $\text{dist}(u)$ is set
 to the distance from s to u .

for all $u \in V$:

$\text{dist}(u) = \infty$

$\text{prev}(u) = \text{nil}$

$\text{dist}(s) = 0$

$H = \text{makequeue}(V)$ (using dist -values as keys)

while H is not empty:

$u = \text{deletemin}(H)$


 for all edges $(u, v) \in E$:

 if $\text{dist}(v) > \text{dist}(u) + l(u, v)$:


$\text{dist}(v) = \text{dist}(u) + l(u, v)$

$\text{prev}(v) = u$


$\text{decreasekey}(H, v)$




A: 0	D: ∞
B: 4	E: ∞
C: 2	




A: 0	D: 6
B: 3	E: 7
C: 2	




A: 0	D: 5
B: 3	E: 6
C: 2	




A: 0	D: 5
B: 3	E: 6
C: 2	




A: 0	D: 5
B: 3	E: 6
C: 2	



- Starting from s , we expand the region R of the graph where shortest paths are known.
- What is the next vertex v to add to R ?
 - the node outside R that is *closest* to s
- Consider the shortest path from s to v . Let u be the node before v on this path.



- Since all edge lengths are positive, u must be closer to s than v is.
- Thus, u is in R . (Since v is the closest node to s outside R .)
- So, the shortest path from s to v is a *known shortest path extended by a single edge*.
- v is the node outside R for which the smallest value of $\text{distance}(s,u) + l(u,v)$ is attained, as u ranges over R .



```

Initialize  $\text{dist}(s)$  to 0, other  $\text{dist}(\cdot)$  values to  $\infty$ 
 $R = \{ \}$  (the ``known region'')
while  $R \neq V$ :
    Pick the node  $v \notin R$  with smallest  $\text{dist}(\cdot)$ 
    Add  $v$  to  $R$ 
    for all edges  $(v, z) \in E$ :
        if  $\text{dist}(z) > \text{dist}(v) + l(v, z)$ :
             $\text{dist}(z) = \text{dist}(v) + l(v, z)$ 

```

Correctness

- Use induction.
- At the end of each iteration of the while loop, the following conditions hold:
 - (1) there is a value d such that all nodes in R are at distance $\leq d$ from s and all nodes outside R are at distance $\geq d$ from s
 - (2) for every node u , the value $\text{dist}(u)$ is the length of the shortest path from s to u whose intermediate nodes are constrained to be in R (if no such path exists, the value is ∞).

```
Initialize  $\text{dist}(s)$  to 0, other  $\text{dist}(\cdot)$  values to  $\infty$ 
 $R = \{ \}$  (the ``known region'')
while  $R \neq V$ :
    Pick the node  $v \notin R$  with smallest  $\text{dist}(\cdot)$ 
    Add  $v$  to  $R$ 
    for all edges  $(v, z) \in E$ :
        if  $\text{dist}(z) > \text{dist}(v) + l(v, z)$ :
             $\text{dist}(z) = \text{dist}(v) + l(v, z)$ 
```

Analysis

- $|V|$ deletemin operations
- $|V| + |E|$ insert/decreasekey operations

Implementation	deletemin	insert/ decreasekey	$ V \times \text{deletemin} + (V + E) \times \text{insert}$
Array	$O(V)$	$O(1)$	$O(V ^2)$
Binary heap	$O(\log V)$	$O(\log V)$	$O((V + E) \log V)$
d -ary heap	$O(\frac{d \log V }{\log d})$	$O(\frac{\log V }{\log d})$	$O((V \cdot d + E) \frac{\log V }{\log d})$
Fibonacci heap	$O(\log V)$	$O(1)$ (amortized)	$O(V \log V + E)$

Update

- We can consider Dijkstra's algorithm as performing a sequence of the following update procedure.

procedure update $((u, v) \in E)$
 $\text{dist}(v) = \min\{\text{dist}(v), \text{dist}(u) + l(u, v)\}$

- This *update operation* uses the fact that the distance to v cannot be more than the distance to $u + l(u, v)$.
- **property 1** It gives the correct distance to v in the particular case where u is the second-last node in the shortest path to v , and $\text{dist}(u)$ is correctly set.
- **property 2** It will never make $\text{dist}(v)$ too small, and in this sense it is *safe*. For instance, extra update's can't hurt.

Property 2 : Initializing $\text{dist}(s) = 0$ and $\text{dist}(v) = \infty$ for all $v \in V - \{s\}$ establishes $\text{dist}(v) \geq \text{distance to } v$ for all $v \in V$.

This invariant is maintained over any sequence of **update**'s.

Pf) Suppose not.

Let v be the first vertex for which $\text{dist}(v) < \text{distance to } v$, and

let u be the vertex that caused $\text{dist}(v)$ to change : $\text{dist}(v) = \text{dist}(u) + l(u, v)$.

$$\begin{aligned} \text{Then, } \text{dist}(v) &< \text{distance to } v && \text{(supposition)} \\ &\leq \text{distance to } u + \text{distance from } u \text{ to } v && \text{(triangle inequality)} \\ &\leq \text{distance to } u + l(u, v) && \text{(shortest path } \leq \text{ specific path)} \\ &\leq \text{dist}(u) + l(u, v) && (v \text{ is first violation}) \end{aligned}$$

Contradiction.

Property 1: Let u be v 's predecessor on a shortest path from s to v . Then, if $\text{dist}(u) = \text{distance to } u$, after **update**(u, v), $\text{dist}(v) = \text{distance to } v$.

Pf) $\text{distance to } v = \text{distance to } u + l(u, v)$.

Suppose $\text{dist}(v) > \text{distance to } v$ before **update**.

(Otherwise, by **property 2**, $\text{dist}(v) = \text{distance to } v$.)

Then, $\text{dist}(v) > \text{dist}(u) + l(u, v)$,


because $\text{dist}(v) > \text{distance to } v = \text{distance to } u + l(u, v) = \text{dist}(u) + l(u, v)$.

By **update**(u, v), $\text{dist}(v) = \text{dist}(u) + l(u, v) = \text{distance to } v$.

Theorem : Dijkstra's algorithm terminates with $\text{dist}(v) = \text{distance to } v$ for all $v \in V$.


Proof.

- It suffices to show that $\text{dist}(v) = \text{distance to } v$ for every $v \in V$ when v is added to R .
- Suppose u is the first vertex added to R for which $\text{dist}(u) > \text{distance to } u$.
- Let y be the first vertex in $V - R$ along a shortest path from s to u , and let x be its predecessor.
- Since u is the first vertex violating the claimed invariant, we have $\text{dist}(x) = \text{distance to } x$.
- When x was added to R , we update the edge (x, y) , which implies that $\text{dist}(y) = \text{distance to } y \leq \text{distance to } u < \text{dist}(u)$.
- But, $\text{dist}(u) \leq \text{dist}(y)$ by our choice of u . Contradiction.




Negative edges

- Dijkstra's algorithm works in part because the shortest path from the starting point s to any node v must pass exclusively through nodes that are closer than v .
- This no longer holds when edge lengths can be negative.



Negative edges

- What is a correct sequence of updates with negative edges?
- Consider a shortest path from s to t :



- This path can have at most $|V|-1$ edges.
- If the sequence of updates performed includes (s, u_1) , (u_1, u_2) , (u_2, u_3) ... (u_k, t) , in that order (*though not necessarily consecutively*), then by **property 1** the distance to t will be correctly computed.
- Simply update *all the edges* $|V|-1$ times!

Bellman-Ford algorithm

procedure shortest-paths (G, l, s)


Input: Directed graph $G = (V, E)$;
 edge lengths $\{l_e : e \in E\}$ with no negative cycles;
 vertex $s \in V$

Output: For all vertices u reachable from s , $\text{dist}(u)$ is set
 to the distance from s to u .

for all $u \in V$:
 $\text{dist}(u) = \infty$
 $\text{prev}(u) = \text{nil}$

$\text{dist}(s) = 0$

repeat $|V| - 1$ times:
 for all $e \in E$:
 update(e)




	Iteration							
Node	0	1	2	3	4	5	6	7
S	0	0	0	0	0	0	0	0
A	∞	10	10	5	5	5	5	5
B	∞	∞	∞	10	6	5	5	5
C	∞	∞	∞	∞	11	7	6	6
D	∞	∞	∞	∞	∞	14	10	9
E	∞	∞	12	8	7	7	7	7
F	∞	∞	9	9	9	9	9	9
G	∞	8	8	8	8	8	8	8

Negative cycles

- If a graph contains a negative-weight cycle, some shortest paths may not exist.
- Instead of stopping after $|V| - 1$ iterations, perform one extra round.
- There is a negative cycle if and only if some dist value is reduced during this final round.

Example:



Shortest paths in dags

- We need to perform a sequence of updates that includes every shortest path as a subsequence.
- In any path of a dag, the vertices appear in increasing linearized order.

procedure dag-shortest-paths(G, l, s)

Input: Dag $G = (V, E)$;

 edge lengths $\{l_e : e \in E\}$; vertex $s \in V$

Output: For all vertices u reachable from s , $\text{dist}(u)$ is set
 to the distance from s to u .

for all $u \in V$:

$\text{dist}(u) = \infty$

$\text{prev}(u) = \text{nil}$

$\text{dist}(s) = 0$

Linearize G

for each $u \in V$, in linearized order:

 for all edges $(u, v) \in E$:

 update(u, v)